

Reinforcement Learning

Policy Gradient, REINFORCE & Actor-Critic Methods

Naeemullah Khan

naeemullah.khan@kaust.edu.sa




جامعة الملك عبد الله
للعلوم والتقنية
King Abdullah University of
Science and Technology



LMH
Lady Margaret Hall
University of Oxford

June 18, 2026



1. Recap
2. Learning Outcomes
3. Motivation for Policy Gradient Methods
4. Policy Gradients
5. REINFORCE
6. Variance Reduction
7. Actor-Critic
8. REINFORCE in action: Recurrent Attention Model (RAM)
9. How to beat the Go world champion - AlphaGo
10. Value-Based and Policy-Based RL
11. Policy Gradient vs. Q-Learning
12. Summary
13. References

- ▶ A Markov Decision Process (MDP) defines the RL problem using:
 - States $\mathcal{S}$
 - Actions $\mathcal{A}$
 - Rewards $\mathcal{R}$
 - Transition probabilities $\mathbb{P}$
 - Discount factor γ
- ▶ Markov property: The future depends only on the present state, not the past.
- ▶ Agent and environment interact in a loop.
- ▶ Policy π decides the agent's actions.

- ▶ Value function: Measures how good a state is.
- ▶ Q-value function: Measures how good a state-action pair is.
- ▶ Bellman equation: Recursively defines value and Q-value functions.
- ▶ Q-learning: Updates Q-values to reduce Bellman error.
- ▶ Deep Q-Learning: Uses neural networks to approximate Q-values.
- ▶ SARSA: On-policy version of Q-learning.

By the end of this session, you will be able to:

- ▶ Understand the **limitations of value-based methods** like Q-learning.
- ▶ Formally define and derive the **Policy Gradient** objective.
- ▶ Implement and interpret the **REINFORCE algorithm**.
- ▶ Explain **Actor-Critic architectures** and their benefits.
- ▶ Evaluate policy-based methods in different environments.

- ▶ Value-based methods learn Q-values and derive the policy indirectly.
 - Inefficient in continuous or large action spaces
 - Can't represent stochastic policies
 - May lead to high variance and instability

Policy Gradient Methods:

- Learn policy parameters directly to maximize expected return:

$$J(\theta) = \mathbb{E}_{\pi_{\theta}}[R]$$



Reinforcement Learning: **Policy Gradients**

- ▶ **What is the problem with Q-learning?**
- ▶ The Q-function can be very complex.
- ▶ For example, a robot grasping an object may have a very high-dimensional state space. It can be difficult to learn the exact Q-value for every (state, action) pair.

- ▶ **What is the problem with Q-learning?**
- ▶ The Q-function can be very complex.
- ▶ For example, a robot grasping an object may have a very high-dimensional state space. It can be difficult to learn the exact Q-value for every (state, action) pair.
- ▶ However, the policy itself can be much simpler; for instance, just closing the robot's hand.
- ▶ Can we learn a policy directly, i.e., find the best policy from a set of possible policies?

- Formally, let us define a class of parameterized policies:

$$\Pi = \{\pi_{\theta} \mid \theta \in \mathbb{R}^m\}$$

- For each policy, we can define its expected return:

$$\mathcal{J}(\theta) = \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r_t \mid \pi_{\theta} \right]$$

- Formally, let us define a class of parameterized policies:

$$\Pi = \{\pi_{\theta} \mid \theta \in \mathbb{R}^m\}$$

- For each policy, we can define its expected return:

$$\mathcal{J}(\theta) = \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r_t \mid \pi_{\theta} \right]$$

- Our goal is to find the optimal policy: $\theta^* = \arg \max_{\theta} \mathcal{J}(\theta)$
- How can we achieve this?


- Formally, let us define a class of parameterized policies:

$$\Pi = \{\pi_{\theta} \mid \theta \in \mathbb{R}^m\}$$

- For each policy, we can define its expected return:

$$\mathcal{J}(\theta) = \mathbb{E} \left[\sum_{t \geq 0} \gamma^t r_t \mid \pi_{\theta} \right]$$

- Our goal is to find the optimal policy: $\theta^* = \arg \max_{\theta} \mathcal{J}(\theta)$
- How can we achieve this?
- **Solution:** Perform gradient ascent on the policy parameters!



Reinforcement Learning: **REINFORCE**

- ▶ REINFORCE is an elegant algorithm for maximizing the expected return.
- ▶ Intuition: trial and error.
- ▶ Sample a trajectory τ . If you get a high reward, try to make it more likely; if you get a low reward, try to make it less likely.
- ▶ A trajectory is a sequence of states, actions, and rewards: $\tau = (s_0, a_0, r_0, s_1, a_1, \dots)$.

- Expected reward:

$$\begin{aligned}\mathcal{J}(\theta) &= \mathbb{E}_{\tau \sim p(\tau; \theta)} [r(\tau)] \\ &= \int_{\tau} r(\tau) p(\tau; \theta) d\tau\end{aligned}$$

- Expected reward:

$$\begin{aligned}\mathcal{J}(\theta) &= \mathbb{E}_{\tau \sim p(\tau; \theta)} [r(\tau)] \\ &= \int_{\tau} r(\tau) p(\tau; \theta) d\tau\end{aligned}$$

- Now let's differentiate this:

$$\nabla_{\theta} \mathcal{J}(\theta) = \int_{\tau} r(\tau) \nabla_{\theta} p(\tau; \theta) d\tau$$

where $p(\tau; \theta) = \prod_{t \geq 0} p(s_{t+1} | s_t, a_t) \pi_{\theta}(a_t | s_t)$

- Expected reward:

$$\begin{aligned}\mathcal{J}(\theta) &= \mathbb{E}_{\tau \sim p(\tau; \theta)} [r(\tau)] \\ &= \int_{\tau} r(\tau) p(\tau; \theta) d\tau\end{aligned}$$

- Now let's differentiate this:

$$\nabla_{\theta} \mathcal{J}(\theta) = \int_{\tau} r(\tau) \nabla_{\theta} p(\tau; \theta) d\tau$$

where $p(\tau; \theta) = \prod_{t \geq 0} p(s_{t+1} | s_t, a_t) \pi_{\theta}(a_t | s_t)$

- But this is **intractable!**

- ▶ However, we can use a useful trick:

$$\begin{aligned}\nabla_{\theta} p(\tau; \theta) &= p(\tau; \theta) \frac{\nabla_{\theta} p(\tau; \theta)}{p(\tau; \theta)} \\ &= p(\tau; \theta) \nabla_{\theta} \log p(\tau; \theta)\end{aligned}$$

- ▶ However, we can use a useful trick:

$$\begin{aligned}\nabla_{\theta} p(\tau; \theta) &= p(\tau; \theta) \frac{\nabla_{\theta} p(\tau; \theta)}{p(\tau; \theta)} \\ &= p(\tau; \theta) \nabla_{\theta} \log p(\tau; \theta)\end{aligned}$$

- ▶ Now, if we substitute this back:

$$\begin{aligned}\nabla_{\theta} \mathcal{J}(\theta) &= \int_{\tau} (r(\tau) \nabla_{\theta} \log p(\tau; \theta)) p(\tau; \theta) d\tau \\ &= \mathbb{E}_{\tau \sim p(\tau; \theta)} [r(\tau) \nabla_{\theta} \log p(\tau; \theta)]\end{aligned}$$

- ▶ We can estimate this with Monte Carlo sampling.

► Recall,

$$p(\tau; \theta) = \prod_{t \geq 0} p(s_{t+1} | s_t, a_t) \pi_{\theta}(a_t | s_t)$$

► Recall,

$$p(\tau; \theta) = \prod_{t \geq 0} p(s_{t+1} | s_t, a_t) \pi_{\theta}(a_t | s_t)$$

► Thus,

$$\log p(\tau; \theta) = \sum_{t \geq 0} \log p(s_{t+1} | s_t, a_t) + \log \pi_{\theta}(a_t | s_t)$$

- Recall,

$$p(\tau; \theta) = \prod_{t \geq 0} p(s_{t+1} | s_t, a_t) \pi_{\theta}(a_t | s_t)$$

- Thus,

$$\log p(\tau; \theta) = \sum_{t \geq 0} \log p(s_{t+1} | s_t, a_t) + \log \pi_{\theta}(a_t | s_t)$$

- When differentiating:

$$\nabla_{\theta} \log p(\tau; \theta) = \sum_{t \geq 0} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

- Recall,

$$p(\tau; \theta) = \prod_{t \geq 0} p(s_{t+1} | s_t, a_t) \pi_{\theta}(a_t | s_t)$$

- Thus,

$$\log p(\tau; \theta) = \sum_{t \geq 0} \log p(s_{t+1} | s_t, a_t) + \log \pi_{\theta}(a_t | s_t)$$

- When differentiating:

$$\nabla_{\theta} \log p(\tau; \theta) = \sum_{t \geq 0} \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

- It doesn't depend on the transition probabilities!

- Therefore, when sampling a trajectory τ , we can estimate $\nabla_{\theta} \mathcal{J}(\theta)$ as:

$$\nabla_{\theta} \mathcal{J}(\theta) \approx \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

- Therefore, when sampling a trajectory τ , we can estimate $\nabla_{\theta} \mathcal{J}(\theta)$ as:

$$\nabla_{\theta} \mathcal{J}(\theta) \approx \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$


- Interpretation:

- If $r(\tau)$ is high, increase the probabilities of the actions taken.
- If $r(\tau)$ is low, decrease the probabilities of the actions taken.


- Therefore, when sampling a trajectory τ , we can estimate $\nabla_{\theta} \mathcal{J}(\theta)$ as:

$$\nabla_{\theta} \mathcal{J}(\theta) \approx \sum_{t \geq 0} r(\tau) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

- Interpretation:
 - If $r(\tau)$ is high, increase the probabilities of the actions taken.
 - If $r(\tau)$ is low, decrease the probabilities of the actions taken.
- It might seem simplistic to say that if a trajectory is good, then all its actions were good. But in expectation, it averages out!



REINFORCE Algorithm Implementation Cycle




Reinforcement Learning: **Variance Reduction**

- ▶ However, there is a problem.
- ▶ This approach suffers from high variance because credit assignment is difficult.
- ▶ Can we help the estimator?

- ▶ However, there is a problem.
- ▶ This approach suffers from high variance because credit assignment is difficult.
- ▶ Can we help the estimator?
- ▶ **First idea:** Increase the probability of an action only by the cumulative future reward from that state:


$$\nabla_{\theta} \mathcal{J}(\theta) \approx \sum_{t \geq 0} \left(\sum_{t' \geq t} r_{t'} \right) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$



- ▶ But this still doesn't completely solve the credit assignment problem.
- ▶ It can lead to bias due to delayed rewards.

- ▶ But this still doesn't completely solve the credit assignment problem.
- ▶ It can lead to bias due to delayed rewards.
- ▶ **Second idea:** Use a discount factor γ to reduce the effect of delayed rewards:

$$\nabla_{\theta} \mathcal{J}(\theta) \approx \sum_{t \geq 0} \left(\sum_{t' \geq t} \gamma^{t'-t} r_{t'} \right) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$



- **Problem:** The raw value of a trajectory isn't necessarily meaningful. For example, if all rewards are positive, you keep increasing the probabilities of actions.

- ▶ **Problem:** The raw value of a trajectory isn't necessarily meaningful. For example, if all rewards are positive, you keep increasing the probabilities of actions.
- ▶ **What is important then?** Whether a reward is better or worse than what you expect to get.

- ▶ **Problem:** The raw value of a trajectory isn't necessarily meaningful. For example, if all rewards are positive, you keep increasing the probabilities of actions.
- ▶ **What is important then?** Whether a reward is better or worse than what you expect to get.
- ▶ **Idea:** Introduce a baseline function dependent on the state:

$$\nabla_{\theta} \mathcal{J}(\theta) \approx \sum_{t \geq 0} \left(\sum_{t' \geq t} \gamma^{t'-t} r_{t'} - b(s_t) \right) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

- ▶ **Problem:** The raw value of a trajectory isn't necessarily meaningful. For example, if all rewards are positive, you keep increasing the probabilities of actions.
- ▶ **What is important then?** Whether a reward is better or worse than what you expect to get.
- ▶ **Idea:** Introduce a baseline function dependent on the state:

$$\nabla_{\theta} \mathcal{J}(\theta) \approx \sum_{t \geq 0} \left(\sum_{t' \geq t} \gamma^{t'-t} r_{t'} - b(s_t) \right) \nabla_{\theta} \log \pi_{\theta}(a_t | s_t)$$

- ▶ A simple baseline: the moving average of rewards experienced so far from all trajectories.

How to Choose the Baseline?

- ▶ Can we choose a better baseline?
- ▶ Essentially, we want to increase the probability of an action from a state if this action was better than the **expected value** from that state.

- ▶ Can we choose a better baseline?
- ▶ Essentially, we want to increase the probability of an action from a state if this action was better than the **expected value** from that state.
- ▶ What does this remind you of?

- ▶ Can we choose a better baseline?
- ▶ Essentially, we want to increase the probability of an action from a state if this action was better than the **expected value** from that state.
- ▶ What does this remind you of?
- ▶ **Answer:** Q-function and value function!

- Intuitively, we are happy with an action a_t in a state s_t if $Q^\pi(s_t, a_t) - V^\pi(s_t)$ is large. In contrast, we are unhappy if it is small.

- ▶ Intuitively, we are happy with an action a_t in a state s_t if $Q^\pi(s_t, a_t) - V^\pi(s_t)$ is large. In contrast, we are unhappy if it is small.
- ▶ The term $Q^\pi(s_t, a_t) - V^\pi(s_t)$ is called the **Advantage** and is denoted by $A^\pi(s_t, a_t)$.

- ▶ Intuitively, we are happy with an action a_t in a state s_t if $Q^\pi(s_t, a_t) - V^\pi(s_t)$ is large. In contrast, we are unhappy if it is small.
- ▶ The term $Q^\pi(s_t, a_t) - V^\pi(s_t)$ is called the **Advantage** and is denoted by $A^\pi(s_t, a_t)$.
- ▶ Using this, we get the estimator:

$$\nabla_\theta \mathcal{J}(\theta) \approx \sum_{t \geq 0} (Q^\pi(s_t, a_t) - V^\pi(s_t)) \nabla_\theta \log \pi_\theta(a_t | s_t)$$



Reinforcement Learning: **Actor-Critic**

- ▶ We do not know the true Q and V functions. Can we learn them?

- ▶ We do not know the true Q and V functions. Can we learn them?
- ▶ Yes, by using Q-learning! We can combine Policy Gradients and Q-learning by training both an **actor** (the policy) and a **critic** (the value or Q-network).

- ▶ We do not know the true Q and V functions. Can we learn them?
- ▶ Yes, by using Q-learning! We can combine Policy Gradients and Q-learning by training both an **actor** (the policy) and a **critic** (the value or Q-network).
- ▶ The actor selects actions, while the critic evaluates how good the chosen actions are and provides feedback for improvement.
- ▶ The two networks adapt to each other, similar to the training process in GANs.

- ▶ We do not know the true Q and V functions. Can we learn them?
- ▶ Yes, by using Q-learning! We can combine Policy Gradients and Q-learning by training both an **actor** (the policy) and a **critic** (the value or Q-network).
- ▶ The actor selects actions, while the critic evaluates how good the chosen actions are and provides feedback for improvement.
- ▶ The two networks adapt to each other, similar to the training process in GANs.
- ▶ This approach simplifies the critic's task, as it only needs to estimate values for (state, action) pairs generated by the current policy.
- ▶ We can also incorporate Q-learning techniques, such as experience replay.

- ▶ We do not know the true Q and V functions. Can we learn them?
- ▶ Yes, by using Q-learning! We can combine Policy Gradients and Q-learning by training both an **actor** (the policy) and a **critic** (the value or Q-network).
- ▶ The actor selects actions, while the critic evaluates how good the chosen actions are and provides feedback for improvement.
- ▶ The two networks adapt to each other, similar to the training process in GANs.
- ▶ This approach simplifies the critic's task, as it only needs to estimate values for (state, action) pairs generated by the current policy.
- ▶ We can also incorporate Q-learning techniques, such as experience replay.
- ▶ **Remark:** The advantage function measures how much better an action was compared to the expected value.

Initialize policy parameters θ , critic parameters ϕ

For iteration=1, 2 ... **do**

 Sample m trajectories under the current policy

$\Delta\theta \leftarrow 0$

For $i=1, \dots, m$ **do**

For $t=1, \dots, T$ **do**

$$A_t = \sum_{t' \geq t} \gamma^{t'-t} r_t^i - V_\phi(s_t^i)$$

$$\Delta\theta \leftarrow \Delta\theta + A_t \nabla_\theta \log(a_t^i | s_t^i)$$

$$\Delta\phi \leftarrow \sum_i \sum_t \nabla_\phi ||A_t^i||^2$$

$$\theta \leftarrow \alpha \Delta\theta$$


$$\phi \leftarrow \beta \Delta\phi$$

End for

REINFORCE in Action: **Recurrent Attention Model (RAM)**

- ▶ **Objective:** Image classification
- ▶ The model takes a sequence of “glimpses,” selectively focusing on regions of the image to predict the class.
 - Inspired by human perception and eye movements
 - Saves computational resources $\Rightarrow$ improves scalability
 - Can ignore clutter or irrelevant parts of the image
- ▶ **State:** Glimpses observed so far
- ▶ **Action:** (x, y) coordinates (center of the next glimpse) indicating where to look next in the image
- ▶ **Reward:** 1 at the final timestep if the image is correctly classified, 0 otherwise


⁰[Mnih et al., 2014]



- ▶ Glimpsing is a non-differentiable operation.
- ▶ The policy for selecting glimpse locations is learned using REINFORCE.
- ▶ Given the sequence of glimpses observed so far, an RNN models the state and outputs the next action.

⁰[Mnih et al., 2014]


REINFORCE in Action: Recurrent Attention Model (RAM)



⁰[Mnih et al., 2014]

How to Beat the Go World Champion: **AlphaGo**

How to Beat the Go World Champion - AlphaGo



- ▶ Combination of supervised learning and reinforcement learning
- ▶ Integration of traditional methods (Monte Carlo Tree Search) with modern approaches (deep reinforcement learning)


⁰[Silver et al., Nature 2016]

- ▶ Featurize the board (stone color, move legality, biases, etc.)
- ▶ Initialize the policy network with supervised training on professional Go games, then continue training using policy gradients (self-play from random previous iterations, with $+1/-1$ reward for winning/losing)
- ▶ Learn a value network (critic) to estimate the value of board positions
- ▶ Finally, combine the policy and value networks within a Monte Carlo Tree Search algorithm to select actions via lookahead search

⁰[Silver et al., Nature 2016]

Reinforcement Learning: **Value-Based** and **Policy-Based** Methods

- ▶ **Value-Based Methods**
 - Learn a value function
 - Derive policy implicitly (e.g., ϵ -greedy)
- ▶ **Policy-Based Methods**
 - Do not learn a value function
 - Learn the policy directly
- ▶ **Actor-Critic Methods**
 - Learn both a value function and a policy



- ▶ Policy gradient and Q-learning use two fundamentally different representations: policies and value functions.
- ▶ Advantage of both methods: no need to model the environment.

- ▶ Policy gradient and Q-learning use two fundamentally different representations: policies and value functions.
- ▶ Advantage of both methods: no need to model the environment.
- ▶ **Policy Gradient: Pros and Cons**
 - **Pros:** Unbiased estimate of the gradient of expected return.
 - Can handle large action spaces (since only one action needs to be sampled).
 - **Cons:** High variance updates (leads to poor sample efficiency).
 - Does not perform credit assignment effectively.

- ▶ Policy gradient and Q-learning use two fundamentally different representations: policies and value functions.
- ▶ Advantage of both methods: no need to model the environment.
- ▶ **Policy Gradient: Pros and Cons**
 - **Pros:** Unbiased estimate of the gradient of expected return.
 - Can handle large action spaces (since only one action needs to be sampled).
 - **Cons:** High variance updates (leads to poor sample efficiency).
 - Does not perform credit assignment effectively.
- ▶ **Q-Learning: Pros and Cons**
 - **Pros:** Lower variance updates, more sample efficient.
 - Performs credit assignment.
 - **Cons:** Biased updates due to function approximation.
 - Difficult to handle large action spaces (since the maximum over actions must be computed).

Reinforcement Learning: **Summary**

Method	Policy Type	Gradient Source	Stability	Efficiency
Q-learning	Deterministic	Value gradients	Medium	High
REINFORCE	Stochastic	Monte Carlo returns	Low	Low
Actor-Critic	Stochastic	TD-based advantage	High	High

► REINFORCE:

- High variance in gradient estimates.
- Slow convergence.

► Actor-Critic:

- Sensitive to hyperparameters.
- Actor and critic updates may interfere.
- Requires careful tuning and exploration strategies.
- Can struggle in sparse-reward environments.

- ▶ **Trust Region Policy Optimization (TRPO):** Improves stability by constraining policy updates.
- ▶ **Proximal Policy Optimization (PPO):** Balances exploration and stability with clipped objective functions.
- ▶ **Soft Actor-Critic (SAC):** Uses entropy regularization for improved robustness and exploration.
- ▶ **Meta-Reinforcement Learning (Meta-RL):** Enables agents to adapt quickly to new tasks.
- ▶ **Multi-agent Actor-Critic:** Facilitates decentralized coordination among multiple agents.

- ▶ It can be hard to learn the exact Q-value for every (state, action) pair in high-dimensional state and action spaces.
- ▶ However, we can just learn a policy that maximizes the reward.
- ▶ We can use gradient ascent on policy parameters.
- ▶ However, this can suffer from high variance. Various strategies exist to tackle this.
- ▶ Actor-Critic methods combine Policy Gradients and Q-learning by training both an actor (the policy) and a critic (the Q-network).
- ▶ The actor decides which action to take, and the critic tells the actor how good its action was and how it should adjust.

Reinforcement Learning: **References**

- [1] Sutton, R. S., & Barto, A. G. (2018).
Reinforcement Learning: An Introduction.

- [2] Williams, R. J. (1992).
Simple Statistical Gradient-Following Algorithms for Connectionist Reinforcement Learning.
Machine Learning Journal.

- [3] Schulman, J., Levine, S., Abbeel, P., Jordan, M., & Moritz, P. (2015).
Trust Region Policy Optimization.
In *ICML*.

- [4] Schulman, J., Wolski, F., Dhariwal, P., Radford, A., & Klimov, O. (2017).
Proximal Policy Optimization Algorithms.
arXiv preprint arXiv:1707.06347.

- [5] Haarnoja, T., Zhou, A., Abbeel, P., & Levine, S. (2018).
Soft Actor-Critic: Off-Policy Maximum Entropy Deep Reinforcement Learning with a
Stochastic Actor.
In *ICML*.

- [6] OpenAI Spinning Up.
<https://spinningup.openai.com>

- [7] Berkeley CS285 Deep RL.
<https://rail.eecs.berkeley.edu/deeprlcourse/>

- [8] Chelsea Finn & Karol Hausman, Stanford CS224R: [Deep Reinforcement Learning](#)
- [9] Fei-Fei Li, Yunzhu Li & Ruohan Gao, Stanford CS231n: [Deep Learning for Computer Vision](#)
- [10] Jimmy Ba & Bo Wang, UofT CSC413/2516: [Neural Networks and Deep Learning](#)
- [11] Sergey Levine, Berkeley CS285: [Deep Reinforcement Learning](#)